


VALUE-BASED APPROACHES


STATE VALUE FUNCTION

$$V_{\pi} : S \rightarrow \mathbb{R}$$

$$V_{\pi}(s) = \mathbb{E} \left[G \mid s(0) = s \wedge \pi \right]$$



BELLMAN EQUATIONS



$$\pi(s_0) = a$$

$$\begin{cases} V_\pi(s_0) = \frac{9}{10}(1 + \gamma V_\pi(s_0)) + \frac{1}{10}(-5 + \gamma V_\pi(s_1)) \\ V_\pi(s_1) = 0 \end{cases}$$

More generally:

$$V_\pi(s) = \mathbb{E}_{s', R} [R + \gamma V_\pi(s') | s, a]$$

optimal state value function:

$$V_* : S \rightarrow \mathbb{R}$$

$$V_*(s) = \sup_{\pi} V_{\pi}(s)$$

greedy policy

$$\pi(s) = \operatorname{argmax}_{a \in A} \mathbb{E} [R + \gamma V(s') \mid s, a]$$

STATE-ACTION VALUE FUNCTION

$$Q_{\pi} : S \times A \rightarrow \mathbb{R}$$

$$Q_{\pi}(s, a) = \mathbb{E} \left[G \mid S(o) = s \wedge A(o) = a \wedge \pi \right]$$

optimal state-action value function:

$$Q_* : S \times A \rightarrow \mathbb{R}$$

$$Q_*(s, a) = \sup_{\pi} Q_{\pi}(s, a)$$

greedy policy

$$\pi(s) = \operatorname{argmax}_{a \in A} Q(s, a)$$

→ FIRST APPROACH : MONTE CARLO

we play the ϵ -greedy policy :

$$\pi(s) = \begin{cases} \operatorname{argmax}_{a \in A} Q(s, a) & \text{with probability } 1 - \epsilon \\ \text{Uniform} & \text{with probability } \epsilon \end{cases}$$

REPEAT

sample a trajectory with ϵ -greedy policy

compute return G_t

$$V(s) \leftarrow V(s) + \alpha (G - V(s))$$

$$\text{NEW} = \text{OLD} + \alpha [\text{CURRENT} - \text{OLD}]$$

→ SECOND APPROACH: TEMPORAL DIFFERENCE

Key idea: bootstrapping

REPEAT

sample a step (s, a, r, s') with ϵ -greedy policy

$$V(s) \leftarrow V(s) + \alpha [r + \gamma V(s') - V(s)]$$

$$\text{NEW} = \text{OLD} + \alpha [\text{CURRENT} - \text{OLD}]$$

Key difference:

- Monte Carlo requires a full trajectory before making an update
BUT smaller variance
- Temporal Difference updates at each step

Q-LEARNING

- Value-based : $Q : S \times A \rightarrow \mathbb{R}$
- Temporal Difference
- Off policy : uses one policy for **acting** and another one for **updating**

REPEAT

sample a step (s, a, r, s') with ϵ -greedy policy

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a' \in A} Q(s', a') - Q(s, a) \right]$$

$$\text{NEW} = \text{OLD} + \alpha [\text{CURRENT} - \text{OLD}]$$

SARSA

- Value-based : $Q : S \times A \rightarrow \mathbb{R}$
- Temporal Difference
- On-policy : uses one policy for acting and for updating

REPEAT

sample a step (s, a, r, s') with ϵ -greedy policy

sample a step (s', a', r', s'') with ϵ -greedy policy

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

$$\text{NEW} = \text{OLD} + \alpha [\text{CURRENT} - \text{OLD}]$$

EXPERIENCE REPLAY

running new samples each time is :

- wasteful
- biased by the current strategy
- not giving the maximum amount of information

EXPERIENCE REPLAY

The buffer has
a fixed size

A buffer stores steps (s, a, r, s')

Two actions:

(1) Sample a trajectory and add each step to the buffer (independently!)

(2) Sample from the buffer to update

PRIORITISED EXPERIENCE REPLAY

- . when we add (s, a, r, s') we compute the bias

$$B = r + \gamma \max_{a' \in A} Q(s', a') - Q(s, a)$$

to obtain a distribution we apply a Softmax:

$$\frac{\exp(B)}{\sum_{B'} \exp(B')}$$

- . we sample from the buffer with this distribution

DOUBLE Q-LEARNING

Issue:

$$Q(s, a) = Q(s, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s', a') - Q(s, a) \right)$$

→ tends to overshoot

THE MAXIMISATION BIAS

X_1, \dots, X_n random variables

Goal: estimate

$$\max_i E[X_i]$$

single estimator:

for each i : Sample X_i a number of times

→ choose $\arg\max_i \hat{E}[X_i] = *$


→ evaluate by $\hat{E}[X_*]$

double estimator:

for each i : Sample X_i a number of times, twice

→ choose $\arg\max_i \hat{E}[X_i] = *$ with first set

→ evaluate by $\hat{E}[X_*]$ with second set



DOUBLE Q-LEARNING ALGORITHM

with probability $1/2$:

- $a' = \arg\max_{a'} Q_2(s', a')$
- $Q_1(s, a) \leftarrow Q_1(s, a) + \alpha \left(r + \gamma Q_1(s', a') - Q_1(s, a) \right)$

else:

symmetrically